

The FLP Theorem

Jacopo Notarstefano

`jacopo.notarstefano [at] gmail.com`

The Distributed Consensus Problem

Definition

Consensus protocol

Definition (Consensus protocol)


A **consensus protocol** is an asynchronous system of N processes ($N \geq 2$). Each process p has a one-bit **input register** x_p , an **output register** y_p with values in $\{\perp, 0, 1\}$ and an unbounded amount of internal storage.

Initial states prescribe fixed starting values for all but the input register; in particular, the output register starts with value \perp .

p acts deterministically according to a **transition** function.

Decision states

The states in which the output register has value 0 or 1 are distinguished as being **decision states**. The transition function cannot change the value of the output register once the process has reached a decision state; that is, the output register is “write-once”.



Message System

Partial correctness

A configuration C has **decision value** v if some process p is in a decision state with $y_p = v$.

Definition (Partial correctness)

A consensus protocol is **partially correct** if:

- 1 No accessible configuration has more than one decision value.
- 2 For each $v \in \{0, 1\}$, some accessible configuration has decision value v .

Total correctness in spite of one fault

A process p is **nonfaulty** in run if it takes infinitely many steps, otherwise it is **faulty**.

A run is **admissible** if at most one process is faulty and all messages sent to nonfaulty processes are eventually received.

A run is **deciding** if some process reaches a decision state.

Definition (Total correctness in spite of one fault)

A consensus protocol P is **totally correct in spite of one fault** if it is partially correct and every admissible run is deciding.

Main result

Theorem (Fischer, Lynch, Paterson 1985)

No consensus protocol is totally correct in spite of one fault.

A configuration is **bivalent** if the set of decision values of configurations reachable from it has 2 elements. It is instead **0-valent** or **1-valent** according to the corresponding value.

Proof (sketch).

Given an initial bivalent configuration, we construct an admissible run that at each stage results in another bivalent configuration. □

Lemma 1


Lemma

Suppose that from some configuration C , the schedules σ_1, σ_2 lead to configurations C_1, C_2 respectively. If the sets of processes taking steps in σ_1 and σ_2 , respectively, are disjoint, then σ_2 can be applied to C_1 and σ_1 can be applied to C_2 , and both lead to the same configuration C_3 .

In other words: **schedules about disjoint processes commute with each other.**

Proof of Lemma 1

Proof (Lemma 1).



Because the sets of processes are disjoint, an event in σ_1 applicable to C is applicable to C_2 as well.

Because of determinism, after all events are processed they must end up in the same state. □

Lemma 2

Lemma

P has a bivalent initial configuration.

Proof (Lemma 2).



Lemma 3

Lemma

Let C be a bivalent configuration of P , and let $e = (p, m)$ be an event that is applicable to C . Let \mathcal{C} be the set of configurations reachable from C without applying e , and let $\mathcal{D} = e(\mathcal{C}) = \{e(E) \mid E \in \mathcal{C} \text{ and } e \text{ is applicable to } E\}$. Then, \mathcal{D} contains a bivalent configuration.

In other words: given a bivalent configuration and an event e applicable to it, **we construct another bivalent configuration having e as the last applied event.**

Proof of Lemma 3

Proof (Lemma 3).



Proof of main result

Proof (main result).

