

# The FLP Theorem

Jacopo Notarstefano

`jacopo.notarstefano [at] gmail.com`

November 26, 2014

# The Distributed Consensus Problem

Let's consider the problem of reaching agreement on a single value in a distributed, asynchronous system of processes.

For example, in a distributed database, all entities that have participated in the processing of a particular transaction need to agree whether to commit or rollback.

Whatever decision is made, all entities must make the same decision in order to preserve the consistency of the database.

# Faults

In this course we usually assumed that the participating processes and the network were completely reliable.

This is not the case with real systems, which are subject to a number of possible **faults**, such as process crashes, network partitioning, and lost, distorted, duplicated messages.

One can even consider more **Byzantine** types of failures, in which faulty processes might go completely haywire, perhaps even sending messages according to some malevolent plan.

# The FLP Theorem

In 1985, Michael Fischer, Nancy Lynch, and Michael Paterson proved the surprising result that no completely asynchronous consensus protocol can tolerate even a single unannounced process death.

Their proof did not consider Byzantine failures, and assumed that the message system is reliable — it delivers all messages correctly and exactly once.

Nevertheless, even with these assumptions, the stopping of a single process at an inopportune time can cause any distributed consensus protocol to fail to reach agreement.

# Consensus protocol


A **consensus protocol** is an asynchronous system of  $N$  processes ( $N \geq 2$ ). Each process  $p$  has a one-bit **input register**  $x_p$ , an **output register**  $y_p$  with values in  $\{\perp, 0, 1\}$  and an unbounded amount of internal storage.

**Initial states** prescribe fixed starting values for all but the input register; in particular, the output register starts with value  $\perp$ .

$p$  acts deterministically according to a **transition** function.

# Decision states

The states in which the output register has value 0 or 1 are distinguished as being **decision states**. The transition function cannot change the value of the output register once the process has reached a decision state; that is, the output register is “write-once”.



# Message system

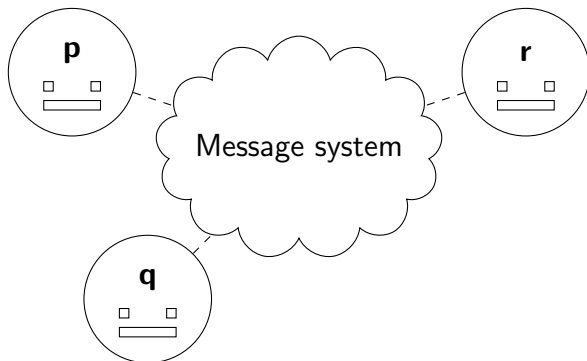
A **message** is a pair  $(p, m)$ , where  $p$  is the name of the destination process and  $m$  is a “message value” from a fixed universe  $M$ .

The **message system** maintains a multiset, called the **message buffer**, of messages that have been sent but not yet delivered. It supports two abstract operations:

- `send(p,m)`: Places  $(p, m)$  in the message buffer.
- `receive(p)`: Deletes some message  $(p, m)$  from the buffer and returns  $m$ , in which case we say  $(p, m)$  is **delivered**, or returns the special null marker  $\emptyset$  and leaves the buffer unchanged.

# Nondeterministic messaging

The message system acts nondeterministically, subject only to the condition that if `receive(p)` is performed infinitely many times, then every message  $(p, m)$  in the message buffer is eventually delivered.





# Configurations and steps

A **configuration** of the system consists of the internal state of each process, together with the contents of the message buffer. An **initial configuration** is one in which each process starts at an initial state and the message buffer is empty.

A **step** takes one configuration to another and consists of a primitive step by a single process  $p$ .

# Primitive steps and applicable events

Let  $C$  be a configuration. A step consists of two phases:

- 1 A  $\text{receive}(p)$  is performed on the message buffer in  $C$  to obtain a value  $m \in M \cup \{\emptyset\}$ .
- 2 Depending on  $p$ 's internal state in  $C$  and on  $m$ ,  $p$  enters a new internal state and sends a finite set of messages to other processes.

The step is completely determined by the pair  $e = (p, m)$ , which we call an **event**.

An event  $e$  that *could happen* at configuration  $C$  is called **applicable**, and the resulting configuration is denoted  $e(C)$ .

# Schedules and runs

A **schedule** from  $C$  is a finite or infinite sequences  $\sigma$  of applicable events, in turn, starting from  $C$ . The associated sequence of steps is called a **run**.

If  $\sigma$  is finite, we let  $\sigma(C)$  denote the resulting configuration, which is said to be **reachable** from  $C$ .

A configuration reachable from some initial configuration is said to be **accessible**.

# Partial correctness

A configuration  $C$  has **decision value**  $v$  if some process  $p$  is in a decision state with  $y_p = v$ .

## Definition (Partial correctness)

A consensus protocol is **partially correct** if:

- 1 No accessible configuration has more than one decision value.
- 2 For each  $v \in \{0, 1\}$ , some accessible configuration has decision value  $v$ .

# Total correctness in spite of one fault

A process  $p$  is **nonfaulty** in run if it takes infinitely many steps, otherwise it is **faulty**.

A run is **admissible** if at most one process is faulty and all messages sent to nonfaulty processes are eventually received.

A run is **deciding** if some process reaches a decision state.

## Definition (Total correctness in spite of one fault)

A consensus protocol  $P$  is **totally correct in spite of one fault** if it is partially correct and every admissible run is deciding.

# Main result

Theorem (Fischer, Lynch, Paterson [FLP85])

*No consensus protocol is totally correct in spite of one fault.*

A configuration is **bivalent** if the set of decision values of configurations reachable from it has 2 elements. It is instead **0-valent** or **1-valent** according to the corresponding value.

Proof (sketch).

Given an initial bivalent configuration, we construct an admissible run that at each stage results in another bivalent configuration. □

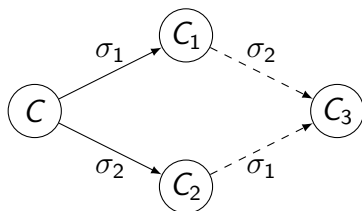
# Lemma 1

## Lemma

*Suppose that from some configuration  $C$ , the schedules  $\sigma_1, \sigma_2$  lead to configurations  $C_1, C_2$  respectively. If the sets of processes taking steps in  $\sigma_1$  and  $\sigma_2$ , respectively, are disjoint, then  $\sigma_2$  can be applied to  $C_1$  and  $\sigma_1$  can be applied to  $C_2$ , and both lead to the same configuration  $C_3$ .*

In other words: **schedules about disjoint processes commute with each other.**

# Proof of Lemma 1



Because the sets of processes are disjoint, an event in  $\sigma_1$  applicable to  $C$  is applicable to  $C_2$  as well.

Because of determinism, after all events are processed they must end up in the same





# Lemma 2

## Lemma

*$P$  has a bivalent initial configuration.*

Assume by contradiction that there isn't. Then, by partial correctness,  $P$  must have both 0-valent and 1-valent initial configurations.

Let's call two initial configurations **adjacent** if they differ only in the initial value  $x_p$  of a single process  $p$ . Any two initial configurations are joined by a chain of initial configurations, each adjacent to the next.

## Proof of Lemma 2

Hence, there must exist a 0-valent initial configuration  $C_0$  adjacent to a 1-valent configuration  $C_1$ . Let  $p$  be the process in whose initial value they differ.

Consider some admissible deciding run from  $C_0$  in which process  $p$  takes no steps, and let  $\sigma$  be the associated schedule. Then  $\sigma$  can be applied to  $C_1$  as well, and the corresponding configurations will be identical except for  $p$ 's internal state.

Both runs reach the same decision value; if it is 1 then  $C_0$  was bivalent, if it is 0 then  $C_1$  was bivalent. □

# Lemma 3

## Lemma

*Let  $C$  be a bivalent configuration of  $P$ , and let  $e = (p, m)$  be an event that is applicable to  $C$ . Let  $\mathcal{C}$  be the set of configurations reachable from  $C$  without applying  $e$ , and let  $\mathcal{D} = e(\mathcal{C}) = \{e(E) \mid E \in \mathcal{C} \text{ and } e \text{ is applicable to } E\}$ . Then,  $\mathcal{D}$  contains a bivalent configuration.*

In other words: given a bivalent configuration and an event  $e$  applicable to it, **we construct another bivalent configuration having  $e$  as the last applied event.**

# Proof of Lemma 3

Proof (Lemma 3).

Assume by contradiction that  $\mathcal{D}$  contains no bivalent configuration. We first show that  $\mathcal{D}$  contains both 0-valent and 1-valent configurations. □

# Proof of main result

Proof (main result).



# Conclusions

TODO

# Bibliography



Michael J. Fischer, Nancy A. Lynch, and Michael S. Paterson, *Impossibility of distributed consensus with one faulty process*, Journal of the ACM (JACM) **32** (1985), no. 2, 374–382.