

# The FLP Theorem

Jacopo Notarstefano

`jacopo.notarstefano [at] gmail.com`

# The Distributed Consensus Problem

## Definition

# Consensus Protocol

# Message System

# Partial correctness

A configuration  $C$  has **decision value**  $v$  if some process  $p$  is in a decision state with  $y_p = v$ .

## Definition (Partial correctness)

A consensus protocol is **partially correct** if:

- 1 No accessible configuration has more than one decision value.
- 2 For each  $v \in \{0, 1\}$ , some accessible configuration has decision value  $v$ .

# Total correctness in spite of one fault

A process  $p$  is **nonfaulty** in run if it takes infinitely many steps, otherwise it is **faulty**.

A run is **admissible** if at most one process is faulty and all messages sent to nonfaulty processes are eventually received.

A run is **deciding** if some process reaches a decision state.

## Definition (Total correctness in spite of one fault)

A consensus protocol  $P$  is **totally correct in spite of one fault** if it is partially correct and every admissible run is deciding.

# Main result

Theorem (Fischer, Lynch, Paterson 1985)

*No consensus protocol is totally correct in spite of one fault.*

A configuration is **bivalent** if the set of decision values of configurations reachable from it has 2 elements. It is instead **0-valent** or **1-valent** according to the corresponding value.

Proof (sketch).

Given an initial bivalent configuration, we construct an admissible run that at each stage results in another bivalent configuration. □

# Lemma 1

## Lemma


*Suppose that from some configuration  $C$ , the schedules  $\sigma_1, \sigma_2$  lead to configurations  $C_1, C_2$  respectively. If the sets of processes taking steps in  $\sigma_1$  and  $\sigma_2$ , respectively, are disjoint, then  $\sigma_2$  can be applied to  $C_1$  and  $\sigma_1$  can be applied to  $C_2$ , and both lead to the same configuration  $C_3$ .*

In other words: **schedules about disjoint processes commute with each other.**



# Proof of Lemma 1

Proof (Lemma 1).



Because the sets of processes are disjoint, an event in  $\sigma_1$  applicable to  $C$  is applicable to  $C_2$  as well.

Because of determinism, after all events are processed they must end up in the same state. □

# Lemma 2

Lemma

*P has a bivalent initial configuration.*

Proof (Lemma 2).



# Lemma 3

## Lemma

*Let  $C$  be a bivalent configuration of  $P$ , and let  $e = (p, m)$  be an event that is applicable to  $C$ . Let  $\mathcal{C}$  be the set of configurations reachable from  $C$  without applying  $e$ , and let  $\mathcal{D} = e(\mathcal{C}) = \{e(E) \mid E \in \mathcal{C} \text{ and } e \text{ is applicable to } E\}$ . Then,  $\mathcal{D}$  contains a bivalent configuration.*

In other words: given a bivalent configuration and an event  $e$  applicable to it, **we construct another bivalent configuration having  $e$  as the last applied event.**

# Proof of Lemma 3

Proof (Lemma 3).



# Proof of main result

Proof (main result).

